

OS Lab Tutorial 2

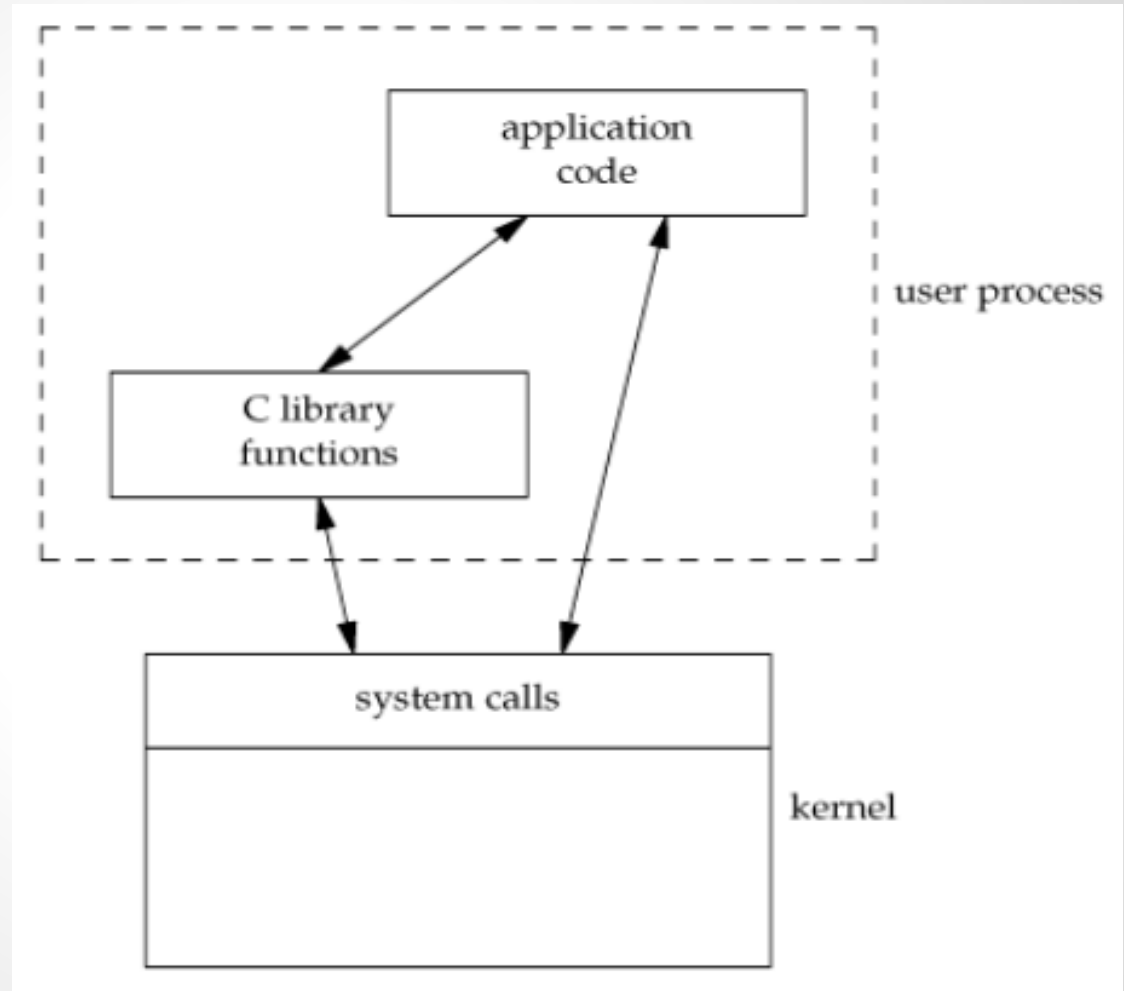
System Calls

System Calls

- Services provided by the kernel
- Documented in Section 2 of the UNIX Programmer's Manual
- Each system call has a function of the same name in the standard C library
- You can use system calls just like the library functions
 - Includes the header files and call the functions
- Examples
 - `man 2 open`
 - `man 2 write`
 - `man 2 fork`

System Calls

- Difference between system calls and library functions



- Reference:

- *Advanced Programming in the UNIX Environment*

Unix Manual

- Section 1: user commands
 - `$ man 1 ls`
 - `$ man 1 printf`
- Section 2: system calls
 - `$ man 2 fork`
 - `$ man 2 getcwd`
- Section 3: general-purpose functions to programmers
 - `$ man 3 printf`
 - `$ man 3 exec`

System Calls - File System

- `mkdir()`: create a new directory
- `rmdir()`: remove a directory
- `getcwd()`: get the current directory
- `chdir()`: change directory
- Use man page to view the details
 - `$ man 2 chdir`
 - `$ man 2 getcwd`

System Calls - Process Management

- **fork()**: create new process
- **getpid()**: return the process identification
- **exec()** - family of functions: execute programs
- **kill()**: send signals to process
- **waitpid()**: wait for process to change state
- Use man page to view the details
 - `$ man 3 exec`
 - `$ man 2 getpid`
 - `exec()` functions (except `execve()`) are actually in the library section. But here we still regard them as system calls

System Calls - fork()

- `#include <unistd.h>`
- After a successful call of `fork()`:
 - A new process is created (child process)
 - Memory space of parent process is copied to child process (So they have the same code and some of the data)
 - child process starts from the call of `fork()`
 - Other details, \$ `man 2 fork`
- One Call, Two Returns
 - Parent gets child's pid (>0), Child gets "0"
 - If the return value is negative, the call fails
 - According to these, we can write corresponding code for child and parent process

System Calls - fork()

- A piece of code in parent process

```
pid_t retValue = fork();           // Call of fork() in parent process

if(retValue == 0)                  // Child gets '0' as the return value of fork()
{
    printf("I'm child process.\n"); // Code for child's behavior
}
else if (retValue > 0)             // Parent gets the pid of child
{
    printf("I'm parent process.\n"); // Code for parent's behavior
}
else                              // Otherwise, it fails to create new proces
{
    printf("Creating Fails.");      // Check errors
}
```

- FYI: Which header file do we need for the type "pid_t" ?

System Calls - exec() functions

- We don't have `exec()`. Instead, we have six functions (`$ man exec`)
- `exec()` are used to load and run a new program so as to replace the current calling process
- `exec()` functions are often used with `fork()` so that child process can execute a new program in its memory space
- If child process doesn't use `exec()`, it is just almost a clone of parent process. `exec()` load new code into the child's memory space

An example of fork() and exec()

- We create a new process using fork and execute a Linux command by one of exec()
- You should use man page to get ideas how to write correct parameters of exec()

```
# include <stdio.h>
```

```
# include <unistd.h>
```

```
int main(void)
```

```
{
```

```
    if (fork() == 0)
```

```
        if (execl("/bin/echo", "echo", "executed by execl", NULL) < 0)
```

```
            perror("Error on execl.");
```

```
    return 0;
```

```
}
```

- **NOTE:** Always check the return value, otherwise, you will lose marks on your code

System Calls - waitpid()

- The calling process will wait for a state change in its child process
- Prototype:
 - header files: "sys/types.h", "sys/wait.h"
 - `pid_t waitpid(pid_t pid, int * status, int options)`
 - e.g., `pid_t retVal = waitpid(child_pid, NULL, WNOHANG);`
- Options
 - `WNOHANG` - Parent process can return immediately if no child has existed
- Exercises
 - Use man page to learn how the three parameters control the behavior of `waitpid()` (E.g., (1) Different pids lead to different actions; (2) What's the meaning of its return values?)
 - Find examples by yourself !

System Calls - kill()

- Send signals to process specified by pid
- prototype
 - header files: "sys/types.h", "signal.h"
 - `int kill(pid_t pid, int signal);`
 - E.g., `int retVal = kill(child_pid, SIGTERM);`
- Exercises
 - Read man page for kill() and learn how to set pid parameter
 - `$ man 2 kill`
 - For more signals, you can read "signal.h" or the man page for signals
 - `$ man 7 signals`
 - What' the difference between **SIGKILL** and **SIGTERM** ?

Background Execution and Process Management

- Switch a program from the foreground to the background or vice-versa.
- Linux Commands
 - **Ctrl + C**: Terminate the foreground process and return to Shell
 - **Ctrl + Z**: Suspend the foreground process, send it to background and return to Shell
 - **jobs**: List background process and their job ID
 - **&**: Let the program run at background
 - **fg [num]**: Move the process with job ID=num to foreground
 - **bg [num]**: Move the process with job ID=num to background
- A nice tutorial [here](#)

System Calls Tool - strace

- Used to trace system calls in your program.

```
/* test.c */  
# include <stdio.h>  
int main(void)  
{  
    printf("Hello World.\n");  
}
```

```
$ gcc -Wall test.c -o test
```

```
$ strace ./test
```

- An interesting and helpful article by Bill Zimmerly

- http://www.ibm.com/developerworks/aix/library/au-unix-strace.html?S_TACT=105AGX52&S_CMP=cn-a-aix

Post Questions in Chat Room